MIN-MARGINAL POWER INDICES FOR MULTI-OBJECTIVE FEATURE SELECTION

Anonymous authors
Paper under double-blind review

ABSTRACT

Game-theoretic feature selection usually aggregates a feature's marginal contributions across coalitions by a mean, which can conceal a context in which adding that feature reverses an acceptable decision. Building on minimum marginal aggregation, we specialize it to a binary utility-threshold voting game and call the resulting quota-indexed support statistic the Min-Marginal Power Index (MinPPI). For a fixed utility quota α, exact MinPPI lies in $\{-1, 0, +1\}$ and distinguishes a feature with a harmful threshold-crossing witness, a feature that is neither harmful nor universally decisive, and a universally decisive feature without requiring a second score cutoff. The minimum analogues of Shapley–Shubik and Banzhaf coincide at the support level because they minimize over the same distinct coalitions. We also study random Monte Carlo and MI-correlation-guided finite-budget estimators, and state the latter's detection result only as a conditional support-mass bound.

We completed a registered seven-dataset campaign, but a post-run integrity audit found that one response was constructed from retained outcome-component predictors; we therefore exclude that dataset from every paper aggregate and report the six leakage-screened datasets over five frozen seeds. Across the resulting 30 dataset–seed units, validation-gated MinPPI removed 12.53% of analyzed features on average, with nonzero removal in 14 units; the descriptive mean held-out ΔU was $-.000313$ across heterogeneous AUC and R^2 scales. This is not a superiority result: MinPPI's mean rank among six removal variants was 4.08 and the Friedman test found no evidence of a method effect ($p = .891$). Under a matched 800-evaluation budget, MinPPI-prefiltered NSGA-II changed dataset-averaged held-out hypervolume by $+0.01584$ (exact paired sign-flip $p = .0625$, Holm-adjusted $p = .4375$), driven mainly by Madelon. Exact-reference estimator evidence on Abalone did not show an advantage for MI-correlation guidance over random sampling. These heterogeneous, low-powered results position MinPPI as a quota-, model-, and split-conditioned candidate generator, not a universally superior feature selector.

1 INTRODUCTION

Feature selection balances predictive utility against the number of retained variables and is therefore naturally multi-objective (Guyon & Elisseeff, 2003; Chandrashekar & Sahin, 2014; Xue et al., 2016). Cooperative-game methods treat features as players and a fitted model's utility as the coalition value, a formulation introduced by Cohen et al. (2007). The Shapley value and Banzhaf index summarize average marginal contributions (Shapley, 1953; Banzhaf, 1965); SHAP made this view widely accessible for model explanation (Lundberg & Lee, 2017). Related selectors include Banzhaf-FS (Karavas et al., 2017), Powershap (Verhaeghe et al., 2022), and ShapG (Zhao et al., 2025). Evolutionary methods such as NSGA-II approximate empirical Pareto fronts but search a combinatorial feature space (Deb et al., 2002; Hamdani et al., 2007).

A mean marginal contribution answers whether a feature helps on average, not whether there exists a coalition in which the feature reverses an otherwise acceptable decision. MinShap replaces averaging by minimum aggregation for general conditional-importance functionals (Zheng & Raskutti, 2026). We specialize that aggregation to a binary utility-threshold game. This pairing supplies

trinary support-level semantics: an observed -1 is existence-certified by an evaluated threshold-crossing coalition and, for a configured quota α, the exact output needs no additional real-valued score cutoff. The quota remains a substantive user choice and is not removed by the construction.

Our contributions are fourfold. First, we define the MinPPI specialization and establish its trinary range, quota-conditioned worst-case semantics, and support-level coincidence of its minimum Shapley–Shubik and minimum Banzhaf analogues. Second, we separate exact computation from finite-budget estimation: observed harmful labels are witness-certified, while missed witnesses remain possible, and the guidance bound is explicitly conditional on the perturbation distribution assigning positive mass to a harmful witness. Third, we use α to generate quota-indexed compression–utility operating points and pass candidate removals through a validation-only gate before test evaluation or downstream search. Fourth, we report a provenance-checked, five-seed analysis on six retained datasets, transparently derived from the completed seven-dataset campaign after a performance-independent data-integrity exclusion. Its evidence is deliberately mixed: removal methods are not statistically separated, prefiltering yields a positive but non-significant average hypervolume change, and guidance behavior depends on dataset, estimator, and utility.

2 METHOD

2.1 UTILITY-THRESHOLD VOTING GAME

Let $N = \{1, \dots, n\}$ denote the analyzed features. For $S \subseteq N$, let $\mathcal{M}_S$ be the fixed model family fitted using S, and let $U(\mathcal{M}_S)$ be a configured real-valued, higher-is-better utility. The empty coalition uses a constant predictor derived only from the training response: class priors for classification and the training mean for regression. With full-feature utility $U_{\text{full}} = U(\mathcal{M}_N) > 0$ and quota $\alpha \in (0, 1]$, define

$$v_\alpha(S) = \mathbb{1}[U(\mathcal{M}_S) \geq \alpha U_{\text{full}}]\,. \tag{1}$$

Unlike a classical monotone simple game, v_α need not be monotone: adding a feature can move utility below the quota. Its marginal is $\Delta_i(S) = v_\alpha(S \cup \{i\}) - v_\alpha(S) \in \{-1, 0, +1\}$.

2.2 MIN-MARGINAL POWER INDICES

For a permutation π of N, let S^i_π be the features preceding i. We define

$$\text{MinSS}_i = \min_\pi \Delta_i(S^i_\pi), \tag{2}$$

$$\text{MinBz}_i = \min_{S \subseteq N \setminus \{i\}} \Delta_i(S). \tag{3}$$

Proposition 1 (exact semantics). For every feature, $\text{MinSS}_i = \text{MinBz}_i \in \{-1, 0, +1\}$. The value is -1 exactly when some coalition is winning before i is added and losing afterward; it is $+1$ exactly when every coalition is losing before i is added and winning afterward; otherwise it is 0. We call these classes *harmful*, *decisive*, and *neutral at the worst-case level*, respectively, always relative to the fixed model family, data split, utility, and quota. The last term means neither harmful nor universally decisive in the exact game; it does not mean that the feature has zero predictive value. The proof is in Appendix A.

Every coalition $S \subseteq N \setminus \{i\}$ occurs as S^i_π for at least one permutation, so the two support-level minima have the same distinct arguments. This does not assert equality of the standard weighted Shapley–Shubik and Banzhaf indices. We write MinPPI_i for either exact minimum analogue.

2.3 FINITE-BUDGET ESTIMATION

Exact evaluation requires at most 2^n distinct coalition utilities and $O(n2^{n-1})$ marginal comparisons. It is practical only at small n. For each feature, **Algorithm A** first samples a coalition size uniformly from $\{0, \dots, n-1\}$ and then samples a subset of that size uniformly without replacement from $N \setminus \{i\}$. For $K \geq 1$, it returns the minimum observed marginal over the frozen schedule $\mathcal{S}_K$:

$$\widehat{\text{MinPPI}}_i = \min_{S \in \mathcal{S}_K} \Delta_i(S) \geq \min_{S \subseteq N \setminus \{i\}} \Delta_i(S). \tag{4}$$

Thus an observed -1 certifies a harmful witness, whereas a sampled 0 or $+1$ does not exclude an unseen smaller marginal.

Algorithm B uses training-only mutual information and absolute correlation to construct a seed coalition: up to ten correlation neighbors are included and then padded from the mutual-information ranking toward $\min(20, \lfloor n/3 \rfloor)$. In each trial, every seed member is removed independently with probability .3; an addition count is drawn uniformly from $\{0, \ldots, \min(5, n-1)\}$ and that many currently absent features are taken without replacement from a random order. The resulting support is restricted by this construction. Conditional on the training data and frozen guidance but before drawing the random schedule, let P_i denote the distribution of one perturbation. If a harmful witness S_i^* has mass $\rho_i = P_i(S_i^*) > 0$ and the $K_{\mathrm{adv}} \geq 1$ draws are independent, then

$$\Pr[\text{detect a harmful witness}] \geq 1 - (1 - \rho_i)^{K_{\mathrm{adv}}}. \tag{5}$$

This geometric statement does not imply that MI-correlation guidance gives larger ρ_i than random sampling. To make the support probabilities in Figure 3(a,b) auditable, consider $N = \{f_0, \ldots, f_4\}$ and the analytically specified game

$$v(S) = \mathbb{1}\big[f_0 \in S \ \wedge \ \big(f_3 \in S \ \vee \ f_4 \in S \ \vee \ (f_2 \in S \wedge f_1 \notin S)\big)\big]. \tag{6}$$

Feature f_1 has the harmful witness $S^* = \{f_0, f_2\}$ and is the only feature with exact minimum marginal -1. Algorithm A samples this witness with probability $p_- = 1/[5\binom{4}{2}] = 1/30$. For illustration only, condition Algorithm B's perturbation kernel on the oracle seed S^* rather than on learned MI-correlation guidance. Its probability of returning exactly S^* is $\rho = .49/5 + 2[.21/(5 \cdot 3)] + .09/(5 \cdot 6) = .129$. Figure 3(b) therefore isolates the perturbation kernel; it is not real-data evidence that learned guidance is better.

2.4 Quota-Indexed Selection and Validation Gate

For each quota, the estimated harmful set $\widehat{\mathcal{H}}_\alpha = \{i : \widehat{\mathrm{MinPPI}}_i = -1\}$ proposes one compression–utility operating point. Sweeping $\alpha \in \{.85, .90, .95, .99\}$ does not solve a Pareto optimization problem and need not produce monotone or nondominated points. The formal removal rule uses the prespecified default $\alpha = .95$.

Every coalition model is fitted on the training split and its game utility is evaluated on validation data. The same validation split is therefore reused adaptively to construct the harmful set, apply the removal gate, and drive NSGA-II search; the held-out split is reserved for registered final evaluation. The candidate subset may trail the full model on validation by at most .005 for classification and .01 for regression. When the raw set fails, the frozen all-or-none backoff removes none of its candidates. The accepted count also defines the deletion budget for SHAP-bottom, MinShap-bottom, Banzhaf-bottom, and random-matched controls. This gate limits validation loss but is not a held-out non-inferiority guarantee.

3 Experimental Design

3.1 Data, Models, and Scope

Table 1 summarizes the six retained datasets. Each seed uses a 64/16/20 train/validation/test split, stratified for classification. Variance filtering, mutual-information selection to at most 100 features, and standardization are fitted on training data only. The frozen seeds are $42, 123, 456, 789, 1024$.

Table 1: Datasets and configured dimension cap. The realized analyzed dimension can be smaller after training-only variance filtering.

Dataset	Native p	Cap	Samples	Task / utility
Spambase	57	57	4,601	binary / ROC AUC
Madelon	500	100	2,600	binary / ROC AUC
ISOLET	617	100	7,797	26-class / OVR AUC
Gisette	5,000	100	7,000	binary / ROC AUC
Credit-Card-Fraud	29	29	284,807	binary / ROC AUC
Abalone	8	8	4,177	regression / R^2

The registered campaign completed seven datasets. A post-run integrity audit found that the Communities & Crime response, `nonViolPerPop`, was within an absolute tolerance of .1 of the sum of four retained outcome-component columns in 2,118 of 2,215 raw rows (95.62%). Those components were retained in every formal seed. We consequently exclude this dataset from all paper aggregates, tests, tables, and figures while preserving the complete original campaign artifacts. This post-hoc exclusion was based on target construction rather than observed method performance; notably, the excluded dataset had contributed a positive held-out hypervolume difference. All reported statistics were recomputed from the canonical row-level records, and the six-dataset derivative is hash-bound to its seven-dataset source.

All retained tasks use XGBoost with probabilities for classification, 100 boosting rounds, depth 4, learning rate .1, one CPU worker per fit, and the CUDA tree backend. The 100-round choice was selected from $\{10, 20, 40, 100\}$ using training/validation-only calibration before formal test evaluation. Binary datasets use ROC AUC, ISOLET uses one-vs-rest AUC, and Abalone uses R^2; ISOLET accuracy appears only in a seed-42 utility ablation.

Removal and full-versus-prefiltered NSGA-II cover all six retained datasets and five seeds. Both NSGA-II variants minimize $|S|/p_0$ and $1-\mathrm{clip}(U, 0, 1)$, where p_0 is the realized analyzed dimension before optional prefiltering and is shared by both variants. They use population 40, 800 evaluations, and reference point $(1.1, 1.1)$. Estimator budgets are $\{2, 5, 10, 25\}$, except ISOLET, which omits 25; exact enumeration is restricted to at most 16 features and therefore provides formal ground truth only on Abalone. The registered α sweep covers the retained datasets at seed 42. Guidance and dispatcher ablations cover Madelon, ISOLET, Credit-Card-Fraud, and Abalone at seed 42; the utility ablation is restricted to ISOLET at seed 42.

3.2 Comparisons, Statistics, and Provenance

The removal comparison includes full, MinPPI, SHAP-bottom, MinShap-bottom, Banzhaf-bottom, and random-matched subsets. Thirty random-matched subset trials are averaged within seed, preserving five paired seed-level replicates. Dataset-level tables report means and sample standard deviations. Across datasets, removal ranks use each method's five-seed dataset mean; a Friedman test and paired Wilcoxon tests with Holm adjustment are descriptive because there are only six blocks.

NSGA-II endpoints are validation hypervolume, held-out hypervolume of the validation-selected front, and a post-hoc attainment count: the first evaluation at which cumulative-best validation hypervolume reaches 95% of that run's final-best value. This diagnostic is neither a stopping time nor attainment of a common quality target. Exact paired sign-flip tests align the five seed replicates and rely on descriptive seed-level sign exchangeability; Holm correction is applied within each seven-comparison endpoint family (six datasets plus the dataset-averaged result). The campaign completed every registered job in the five-seed core stages and the seed-42 sensitivity stages, with no approved omissions or skipped jobs. The original seven-dataset summary remains immutable; the integrity-screened tables, tests, and figures were rederived and independently hash-verified before the paper was updated.

4 Results

4.1 RQ1: Validation-Gated Removal

Table 2 shows heterogeneous compression. Averaged over all 30 retained dataset–seed units, MinPPI removed 12.53% of analyzed features; the descriptive mean held-out ΔU was $-.000313$ across heterogeneous AUC and R^2 scales, and 14 of 30 units removed at least one feature. Madelon had the largest mean compression (43.6%) and a positive mean $\Delta U = +.00469$, whereas Abalone compressed by 27.5% with mean $\Delta U = -.00565$; Credit-Card-Fraud compressed by 2.8% with held-out $\Delta U = -0.001$. The raw harmful set passed its validation gate in 27 of 30 units, including 13 vacuous passes with no proposed removal; among the 17 nonempty candidate sets, 14 passed. The failed cases—Abalone seeds 42 and 1024 and Madelon seed 1024—backed off to zero removal, demonstrating that the gate is a substantive part of the procedure.

Table 2: MinPPI removal at $\alpha = .95$ over five seeds. Compression and held-out $\Delta U = U_{\text{selected}} - U_{\text{full}}$ are mean $\pm$ sample SD. "Raw pass" counts validation-gate passes before backoff, including vacuous passes when no removal was proposed.

Dataset	Compression (%)	Held-out ΔU	Raw pass
Abalone	27.50 ± 25.62	$-.00565 \pm .01077$	$3/5$
Credit-Card-Fraud	2.76 ± 2.89	$-.00078 \pm .00335$	$5/5$
Gisette	$.60 \pm .89$	$-.00011 \pm .00016$	$5/5$
ISOLET	$.00 \pm .00$	$.00000 \pm .00000$	$5/5$
Madelon	43.60 ± 31.96	$+.00469 \pm .00497$	$4/5$
Spambase	$.72 \pm .99$	$-.00003 \pm .00007$	$5/5$

The descriptive, unweighted mean of the six dataset-level utilities was .877807 for MinPPI and .878121 for the full model, a difference of $-.000313$ (Wilcoxon $p = .4375$, Holm $p = 1$); because ROC AUC, OVR AUC, and R^2 are mixed, this mean has no single common utility unit. MinPPI's mean rank among all six variants was 4.08. The omnibus Friedman test found no evidence of a method effect ($\chi^2 = 1.686$, $p = .891$), and every MinPPI pairwise Holm-adjusted p-value was 1. Figure 1 therefore visualizes effect heterogeneity rather than a critical-difference or superiority result. Random matched deletion was particularly unfavorable on Madelon ($\Delta U = -.0732$), but the formal ranking does not establish MinPPI as globally best.

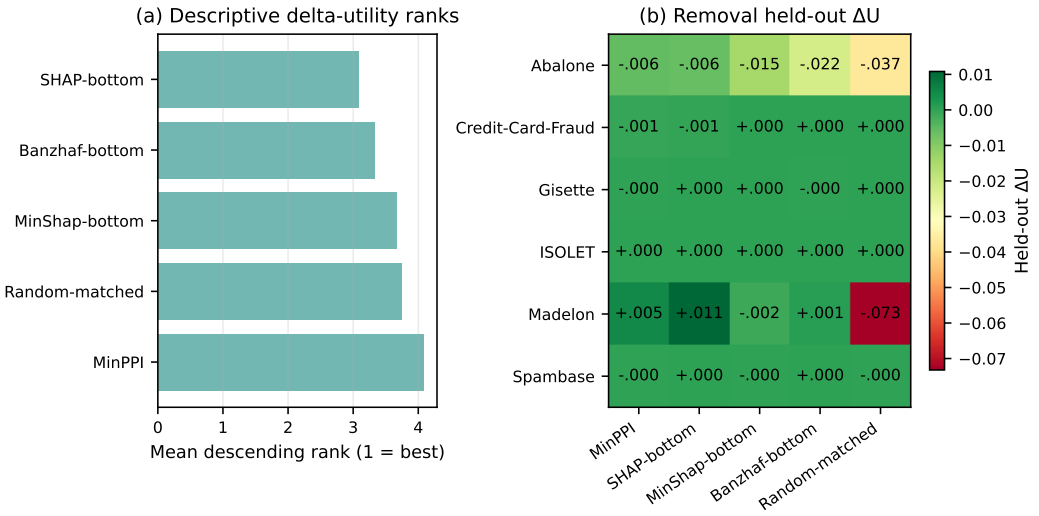


Figure 1: Validation-gated matched removal. (a) Descriptive mean descending ranks (1 is best) among all six methods across six retained datasets, computed from each method's five-seed mean ΔU; the full model participates in the ranking but is omitted from the display. (b) Five-seed mean ΔU, where positive values favor removal. Each control uses the accepted MinPPI-matched removal count, and 30 random-matched subset trials are averaged within seed. The display is descriptive and does not establish a broad method ordering.

4.2 RQ2: Quota Sensitivity

The α sweep is a seed-42 sensitivity analysis ($n = 1$ per point), not a model-selection exercise. Its trajectories were dataset-dependent and non-monotone. On Madelon, compression across $\alpha = (.85, .90, .95, .99)$ was $(.24, .49, .56, .51)$ with held-out $\Delta U = (+.0077, +.0227, -.0003, -.0054)$. On Abalone, compression was $(.375, .50, .50, 0)$, while on Credit-Card-Fraud it rose from 0 at .85 and .90 to .034 at .95 and .586 at .99. Figure 3(c) overlays the registered points for Madelon and Credit-Card-Fraud on held-out evaluations of their seed-42 validation-selected fronts. The points illustrate quota sensitivity only; they do not prove monotonicity, held-out nondominance, or cross-seed safety.

4.3 RQ3: Prefiltered Multi-Objective Search

With the same 800-evaluation budget, MinPPI-prefiltered NSGA-II changed the dataset-averaged held-out hypervolume by $+.01584$ and validation hypervolume by $+.01741$. For each endpoint, all five seed-level pooled differences were positive (rank-biserial correlation 1), but the exact sign-flip $p = .0625$ became $.4375$ after the registered Holm correction. Dataset mean test-HV differences were $+.09515$ on Madelon, $+.00535$ on Credit-Card-Fraud, $+.00021$ on Abalone, 0 on ISOLET, $-.00191$ on Gisette, and $-.00378$ on Spambase. Thus the positive pooled direction is driven mainly by Madelon, with only small positive differences on Credit-Card-Fraud and Abalone.

The post-hoc 95%-of-final-HV attainment diagnostic occurred 22.67 evaluations earlier on average ($p = .4375$, Holm $p = 1$), with the largest decreases on Madelon and Gisette and an increase on Spambase. Both variants nevertheless consumed the same registered total budget of 800 evaluations. Figure 2 therefore supports a directional, dataset-dependent change in search outcomes, not universal improvement, budget reduction, or runtime acceleration.

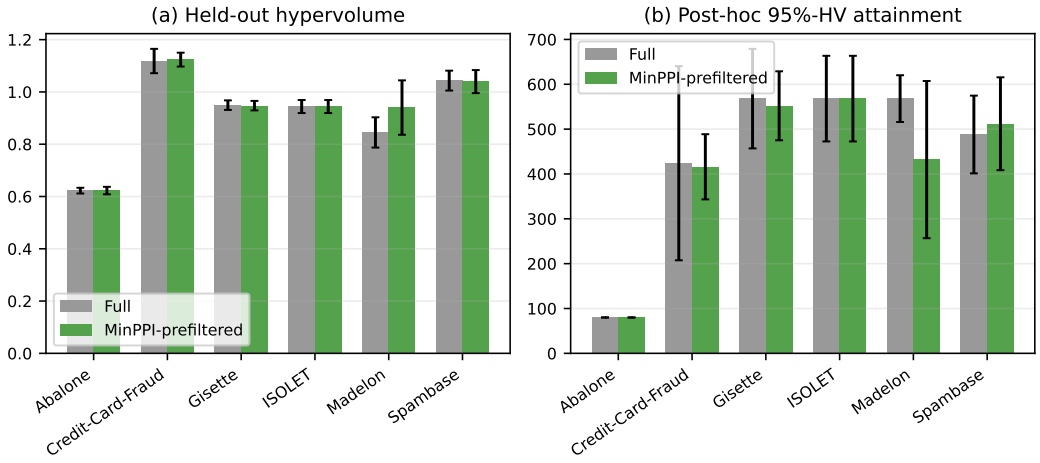


Figure 2: Full versus validation-gated MinPPI-prefiltered NSGA-II under the matched 800-evaluation budget. Bars are five-seed means and whiskers are sample standard deviations for (a) held-out hypervolume of the validation-selected front, using reference point $(1.1, 1.1)$, where the dataset-averaged held-out change is +0.016, and (b) the post-hoc first evaluation attaining 95% of that run's final-best validation hypervolume, where lower is earlier. The latter is not a stopping time. No dataset-wise comparison survives the registered exact paired sign-flip tests with within-endpoint Holm correction.

4.4 RQ4: Estimator, Guidance, and Utility Sensitivity

Abalone is the only retained dataset small enough for exact enumeration. Averaged over the five seed-level F1 scores, at budgets $2, 5, 10, 25$, random MC achieved $.100, .150, .230, .443$, while MI-correlation guidance achieved $.000, .080, .133, .313$. The hybrid dispatcher scored 1 at every budget because it selected exact enumeration at eight features; this does not validate the approximate guidance heuristic. Thus the five-seed mean F1 for learned guidance was lower than random sampling at every registered finite budget; this is not a claim about every individual seed.

The seed-42 ablations likewise show sensitivity rather than a stable guidance ordering. On Madelon, random, MI-only, correlation-only, and MI-correlation seeds recovered 5, 16, 34, and 28 harmful labels; on Credit-Card-Fraud the counts were 6, 2, 2, and 1. On ISOLET, changing the utility while keeping the probability-producing classifier fixed changed the harmful count from 0 under the formal OVR AUC to 9 under accuracy. These single-seed results show that estimator and utility choices materially affect recovered labels, but they cannot support significance or cross-seed stability claims.

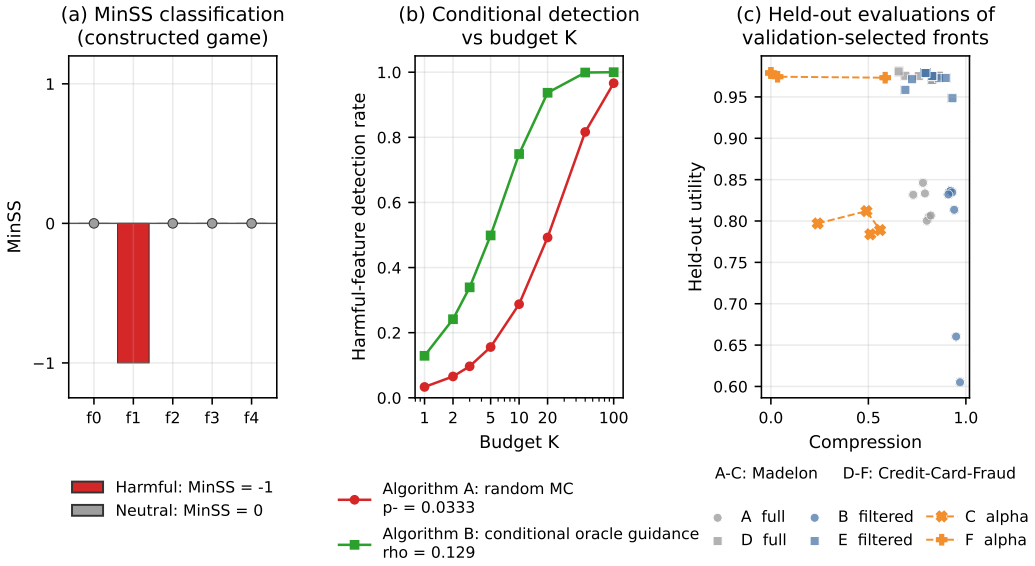


Figure 3: Overview of MinPPI. (a) Exact minimum marginals in the five-feature game of equation 6: $f_1 = -1$ and all other features are neutral at the worst-case level. (b) Geometric detection probabilities for the unique harmful witness under random MC ($p_- = 1/30 = .0333$) and Algorithm B's perturbation kernel conditional on the oracle seed $\{f_0, f_2\}$ ($\rho = .129$); this isolates the kernel and is not real-data guidance evidence. (c) Held-out evaluations of seed-42 validation-selected fronts for full and gated prefiltered NSGA-II, together with the registered $\alpha \in \{.85, .90, .95, .99\}$ operating points on Madelon and Credit-Card-Fraud. Connecting lines only order the α values and do not assert held-out nondominance.

5 Discussion

The exact contribution of the MinPPI specialization is sharper than its empirical claim. Given α, an exact or sampled -1 is existence-certified by an evaluated threshold-crossing coalition and needs no additional score threshold; the validation gate then separates witness discovery from the decision to remove a feature. Finite sampling preserves witness validity but not witness completeness, and sampled 0 or $+1$ does not certify the corresponding exact class. This distinction is important because the retained study does not support a universal selector ranking: MinPPI stayed close to full-model utility descriptively, yet had mean rank 4.08 among six methods, and Abalone exhibited a negative mean held-out effect.

The downstream results are similarly conditional. Madelon accounts for most of the positive average hypervolume change; Credit-Card-Fraud and Abalone show only small positive differences, while Gisette and Spambase show small decreases. This pattern is compatible with prefiltering being useful when its accepted set removes distracting dimensions, but the study neither identifies a general dataset condition nor establishes statistical superiority. The registered validation gate also cannot guarantee held-out non-inferiority; it uses no test information by design.

Several limitations constrain interpretation. First, five paired seed replicates and six retained datasets give low-powered tests; seed-level sign-flip inference is descriptive and depends on sign exchangeability. The α, guidance, dispatcher, and utility ablations use only seed 42. The necessary post-hoc integrity exclusion also means that the reported analysis set was not fully prespecified, despite the performance-independent exclusion rule and preservation of the original records; independent replication on prospectively screened data is needed. Second, the study uses one XGBoost family and utilities on different scales; a common multiplicative quota is not invariant to shifts of the utility scale. Third, exact estimator comparison is available only on eight-feature Abalone, where the five-seed mean F1 for guidance was lower than random sampling at every registered finite budget. The

bound in equation 5 is a support condition, not evidence that the guidance distribution assigns more mass to harmful witnesses. Fourth, shared caching and execution order complicate runtime attribution, so we make no computational speedup claim. Finally, the current figures summarize means and sample standard deviations; paired seed-level variation and broader model families remain important future work.

6 CONCLUSION

Exact MinPPI specializes minimum marginal aggregation to a binary utility-threshold game and yields a quota-conditioned three-way classification; finite-budget estimators certify only observed harmful threshold crossings. The resulting candidates feed a validation-gated removal or prefiltering step. Across the six integrity-screened datasets, the method produced nonzero and sometimes substantial compression on some datasets and a positive descriptive average NSGA-II hypervolume change, but neither removal superiority nor a general guidance advantage was detected. The appropriate conclusion is therefore methodological and conditional: MinPPI is a transparent candidate generator whose value depends on the model, split, dataset, estimator, utility, quota, and validation decision.

REFERENCES

John F. Banzhaf. Weighted voting doesn't work: A mathematical analysis. *Rutgers Law Review*, 19: 317–343, 1965.

Girish Chandrashekar and Ferat Sahin. A survey on feature selection methods. *Computers & Electrical Engineering*, 40(1):16–28, 2014.

Shay Cohen, Gideon Dror, and Eytan Ruppin. Feature selection via coalitional game theory. *Neural Computation*, 19(7):1939–1961, 2007.

Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and T. Meyarivan. A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evolutionary Computation*, 6(2):182–197, 2002.

Isabelle Guyon and André Elisseeff. An introduction to variable and feature selection. *Journal of Machine Learning Research*, 3:1157–1182, 2003.

Tarek M. Hamdani, Jin-Myung Won, Adel M. Alimi, and Fakhri Karray. Multi-objective feature selection with NSGA II. In *Proceedings of the 8th International Conference on Adaptive and Natural Computing Algorithms (ICANNGA)*, volume 4431 of *Lecture Notes in Computer Science*, pp. 240–247. Springer, 2007.

Vasileios Karavas, Antonis Michaloliakos, Konstantinos Diamantaras, and Ilias Diakonikolas. Feature importance scores and lossless feature pruning using Banzhaf power indices. *arXiv preprint arXiv:1711.04992*, 2017.

Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30, 2017.

Lloyd S. Shapley. A value for n-person games. In Harold W. Kuhn and Albert W. Tucker (eds.), *Contributions to the Theory of Games*, volume 2, pp. 307–317. Princeton University Press, 1953.

Jarne Verhaeghe, Jonas Van Der Donckt, Femke Ongenae, and Sofie Van Hoecke. Powershap: A power-full Shapley feature selection method. In *Proceedings of the European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML PKDD)*, 2022.

Bing Xue, Mengjie Zhang, Will N. Browne, and Xin Yao. A survey on evolutionary computation approaches to feature selection. *IEEE Transactions on Evolutionary Computation*, 20(4):606–626, 2016.

Chi Zhao, Jiapeng Liu, and Elena Parilina. ShapG: New feature importance method based on the Shapley value. *Engineering Applications of Artificial Intelligence*, 148, 2025.

Chenghui Zheng and Garvesh Raskutti. MinShap: A Shapley-based framework for feature redundancy. *arXiv preprint arXiv:2604.15107*, 2026. Version 2, revised July 2026.

A PROOFS

Proof of Proposition 1. Because $v_\alpha \in \{0, 1\}$, every marginal $\Delta_i(S)$ is in $\{-1, 0, +1\}$, and so is any nonempty minimum of such marginals. For each $S \subseteq N \setminus \{i\}$ there are $|S|!(n - |S| - 1)! \geq 1$ permutations for which S is exactly the predecessor set of i. Hence the distinct arguments of MinSS_i and MinBz_i coincide; multiplicity cannot change a minimum.

The minimum equals -1 if and only if at least one marginal is -1, which is equivalent to the existence of S with $v_\alpha(S) = 1$ and $v_\alpha(S \cup \{i\}) = 0$. It equals $+1$ if and only if every marginal is $+1$, which is equivalent to $v_\alpha(S) = 0$ and $v_\alpha(S \cup \{i\}) = 1$ for every S. If neither condition holds, the trinary range forces the minimum to be 0. □

Finite-sample one-sidedness. For any sampled collection $\mathcal{S}_K \subseteq 2^{N \setminus \{i\}}$,

$$\min_{S \in \mathcal{S}_K} \Delta_i(S) \geq \min_{S \subseteq N \setminus \{i\}} \Delta_i(S),$$

because the left minimum is taken over a subset of the right minimum's arguments. An observed value -1 is therefore exact evidence of at least one harmful witness, but a larger sampled minimum can result from an unseen smaller marginal. □

Conditional detection bound. Let independent perturbations be drawn from P_i, and suppose a fixed harmful witness S_i^* has mass $\rho_i > 0$. One draw misses that witness with probability $1 - \rho_i$; all K_{adv} draws miss it with probability $(1-\rho_i)^{K_{\mathrm{adv}}}$. Sampling S_i^* implies detection, which yields equation 5. The statement is silent when $\rho_i = 0$ and does not compare P_i with the random-MC distribution. □

B ADDITIONAL FORMAL RESULTS

B.1 QUOTA SWEEP

Table 3 reports the seed-42 quota sweep for the six retained datasets. Compression is the accepted removed fraction and ΔU is held-out selected-minus-full utility. The non-monotone rows show why $\alpha = .95$ was fixed before the campaign rather than chosen retrospectively.

Table 3: Seed-42 quota operating points, shown as (compression, held-out ΔU). Every cell has $n = 1$.

Dataset	$\alpha = .85$	$\alpha = .90$	$\alpha = .95$	$\alpha = .99$
Abalone	$(.375, +.0041)$	$(.500, -.0198)$	$(.500, +.0010)$	$(0, 0)$
Credit-Card-Fraud	$(0, 0)$	$(0, 0)$	$(.034, -.0045)$	$(.586, -.0057)$
Gisette	$(0, 0)$	$(0, 0)$	$(0, 0)$	$(.090, +.0002)$
ISOLET	$(0, 0)$	$(0, 0)$	$(0, 0)$	$(.010, -.000005)$
Madelon	$(.240, +.0077)$	$(.490, +.0227)$	$(.560, -.0003)$	$(.510, -.0054)$
Spambase	$(0, 0)$	$(0, 0)$	$(0, 0)$	$(.196, -.0012)$

B.2 ESTIMATOR AND SEED-42 ABLATIONS

Table 4: Harmful-label F1 against exact Abalone labels over five seeds. Exact and hybrid equal 1 at all budgets because the hybrid dispatches to exact enumeration at eight features.

Estimator	$K = 2$	$K = 5$	$K = 10$	$K = 25$
Random MC	.100	.150	.230	.443
MI-correlation guided	.000	.080	.133	.313
Exact	1.000	1.000	1.000	1.000
Hybrid dispatcher	1.000	1.000	1.000	1.000

Table 5: Seed-42 ablation harmful counts. A dash means exact enumeration was disallowed by the registered dimension limit, not that zero harmful features were found.

Axis	Variant	Abalone	Credit	ISOLET	Madelon
Guidance	random	2	6	1	5
	MI only	2	2	0	16
	correlation only	2	2	0	34
	MI-correlation	2	1	0	28
Dispatcher	exact	4	–	–	–
	random MC	2	2	0	34
	MI-correlation	2	1	0	28
	hybrid	4	1	0	56

For ISOLET, the registered utility ablation found 0 harmful features with formal OVR AUC and 9 with accuracy. The classifier family and probability outputs were held fixed; the difference is therefore attributable to the utility definition within this seed-42 comparison.

C Reproducibility Details

The formal splits contain 64% training, 16% validation, and 20% test observations. Dataset identities, split indices, transformation metadata, model policy, source identity, and the scientific software environment are recorded with cryptographic digests. Resume operations reject protocol drift. Training data fit preprocessing, guidance, and every coalition model. Validation utility is reused adaptively for threshold-game construction, the removal gate, and NSGA-II selection. The held-out partition is not used to choose features, the removal backoff, α, or NSGA-II individuals; it is read only for registered outcome evaluation after validation selection. The complete seven-dataset summary remains the historical record; the six-dataset publication view records the exclusion audit, source digest, regenerated tables, and regenerated figure digests.

The formal estimator set comprises exact enumeration where permitted, random MC, MI-correlation guidance, and the hybrid dispatcher. Removal controls delete the accepted MinPPI-matched count. NSGA-II-full and NSGA-II-prefilter use identical population, variation, evaluation budget, reference point, and validation-selection logic. Runtime, cache-hit, and physical-fit counters are diagnostic only and are excluded from predictive superiority claims.