MIN-MARGINAL POWER INDICES FOR MULTI-OBJECTIVE FEATURE SELECTION: A VOTING-GAME APPROACH

Anonymous authors
Paper under double-blind review

ABSTRACT

Existing game-theoretic feature selection methods mostly score each feature by its *mean* marginal contribution across coalitions. The mean, however, conceals a critical risk: a feature that is beneficial on average may be actively harmful in specific feature combinations, and no existing method systematically identifies such features before they degrade deployed model performance. We address this gap by introducing Min-Marginal Power Indices (MinPPI). Rather than averaging, MinPPI takes the *minimum* marginal contribution of each feature under a binary voting game $v_\alpha(S) = \mathbb{1}[\text{Acc}(\mathcal{M}_S) \geq \alpha \cdot \text{Acc}_{\text{full}}]$. This min–binary pairing yields a discrete output in $\{-1, 0, +1\}$, cleanly classifying every feature as *harmful* (must remove), *neutral*, or *decisive* (must retain). We prove the equivalence of the min Shapley-Shubik and min Banzhaf formulations, and develop a ShapG-guided adversarial search that achieves a better sample efficiency gain over uniform Monte Carlo for detecting sparse harmful features. Embedding MinPPI into the mutation and crossover operators of NSGA-II yields the MOFS-MinPPI algorithm, with a provable safety guarantee: MinPPI estimation errors affect convergence speed but never compromise Pareto-optimality. On five tabular datasets (29–5,000 features), MOFS-MinPPI significantly outperforms seven baselines—SHAP, ShapG, MinShap, Banzhaf, Powershap, MOFS-SHAP, and NSGA2-uniform—in Pareto-front hypervolume and compression rate at 1%/3%/5% AUC-loss thresholds ($p < 0.05$, Friedman–Nemenyi). Ablation confirms that the min operator consistently dominates the mean and that binary v_α dominates continuous $v(\cdot)$ across all conditions. These findings suggest that replacing the mean with the minimum over a thresholded value function offers a broadly applicable principle for robust feature selection beyond the specific methods studied here.

1 INTRODUCTION

Feature selection—identifying a compact yet predictive subset from hundreds or thousands of candidate features—is a foundational preprocessing step that directly governs model interpretability, training cost, and generalization (Guyon & Elisseeff, 2003; Chandrashekar & Sahin, 2014). The underlying problem is combinatorial: the space of possible subsets grows exponentially in dimensionality, rendering exhaustive search intractable beyond trivial feature counts.

Real-world feature selection is inherently multi-objective. Practitioners balance at least three competing criteria: minimizing cardinality (interpretability and deployment cost), maximizing predictive accuracy, and minimizing training time (experimentation velocity). Single-objective approaches—filter methods based on univariate statistics, or wrapper methods optimizing scalarized objectives—compress this trade-off into a single number, concealing the full solution landscape (Jiao et al., 2024). Multi-objective evolutionary algorithms, notably the Non-dominated Sorting Genetic Algorithm II (NSGA-II) (Deb et al., 2002), recover the complete Pareto front of non-dominated trade-offs and have been adapted successfully to feature selection (Hamdani et al., 2007; Song et al., 2024; Shao et al., 2025).

Cooperative game theory offers a complementary, axiomatically grounded perspective. Casting each feature as a player in a coalitional game where coalition value is the predictive performance attain-

able with those features alone, the Shapley value (Shapley, 1953), Shapley-Shubik index (Shapley & Shubik, 1954), and Banzhaf index (Banzhaf, 1965) quantify each feature's marginal contribution across all possible coalitions. This framework entered mainstream machine learning through SHAP (Lundberg & Lee, 2017). A critical limitation, however, is that classical indices aggregate marginal contributions via the *mean* over permutations or subsets. The mean smooths contextual diversity: a feature beneficial in most contexts but actively harmful in a few receives a misleadingly positive score, concealing the harmful contexts entirely. Recent game-theoretic feature selection methods—Banzhaf-based pruning (Karavas et al., 2017), Powershap (Verhaeghe et al., 2022), and ShapG (Zhao et al., 2025)—all inherit this limitation through mean-based aggregation.

MinShap (Zheng & Raskutti, 2026) recently proposed a conceptually important shift: replacing the mean with the *minimum* over permutations. By capturing the worst-case marginal, MinShap exposes features that are detrimental in specific contexts—information the mean discards. We regard MinShap as inspiration; our theoretical development is independent. Importantly, MinShap operates on a continuous value function $v(S) = \text{AUC}(\mathcal{M}_S)$, so its output remains a continuous quantity in $[-1, 1]$, lacking a principled decision boundary. We argue that a more structured framework emerges from a *binary* voting-game formulation (Owen, 1972): $v_\alpha(S) = \mathbb{1}[\text{Acc}(\mathcal{M}_S) \geq \alpha \cdot \text{Acc}_{\text{full}}] \in \{0, 1\}$, under which a coalition either meets a fraction α of full-model accuracy or does not. Pairing this binary v_α with the min operator yields a discrete three-valued output in $\{-1, 0, +1\}$, cleanly classifying features as *harmful*, *neutral*, or *decisive*.

This formulation also sidesteps a known structural defect. Direct WVG-FS methods approximate coalition value through an additive surrogate $v(S) \approx \mathbb{1}[\sum_i w_i \geq q]$ (Cohen et al., 2007), which is systematically inaccurate for non-linear models. In our framework, $v_\alpha(S)$ is evaluated by training $\mathcal{M}_S$ and measuring accuracy—no additive surrogate is employed. The MinPPI scores serve exclusively as soft guidance for NSGA-II variation operators; final Pareto-optimality is determined from actual model fitness (Section 4.5, Proposition 2).

We synthesize three independent lines of investigation—game-theoretic power indices, the min-marginal operator, and multi-objective evolutionary search—into a unified feature selection framework. Our contributions are:

1. **Minimal Power Indices (MinPPI).** We formalize the Min Shapley-Shubik (MinSS) and Min Banzhaf (MinBz) indices under v_α. We prove $\text{MinSS}_i = \text{MinBz}_i$ (Lemma 2) and that both take values in $\{-1, 0, +1\}$ (Lemma 1), yielding a three-way classification: decisive $(+1)$, harmful (-1), neutral (0) (Theorem 1).

2. **ShapG-Guided Adversarial Search (Algorithm B).** Random Monte Carlo for detecting harmful features suffers $O(1/p_-)$ sample complexity when harmful contexts are sparse. We leverage the ShapG correlation graph (Zhao et al., 2025) to guide adversarial search toward vulnerable coalitions, achieving a $\mathbf{38}\times$ sample efficiency gain in detecting sparse harmful features.

3. **MOFS-MinPPI Algorithm.** We embed MinPPI into NSGA-II operators: harmful features receive elevated mutation probability biased toward removal; decisive features are protected. Guidance strength is controlled by a single parameter β, with $\beta = 0$ recovering the unguided baseline.

4. **Comprehensive Empirical Validation.** We evaluate MinPPI-MOFS against seven baselines on five tabular datasets (29–5,000 features, five seeds each). Performance is measured via Pareto-front hypervolume (HV), compression rate at $1\%/3\%/5\%$ AUC-loss thresholds, and training-time speedup. Ablation isolates the min vs. mean operator, binary vs. continuous $v(\cdot)$, and guidance strength β, with statistical significance via the Friedman–Nemenyi test.

2 Related Work

We survey three intersecting threads: game-theoretic feature importance, multi-objective feature selection, and voting games in machine learning.

2.1 GAME-THEORETIC FEATURE IMPORTANCE

The Shapley value (Shapley, 1953) is the unique attribution scheme satisfying linearity, dummy, symmetry, and efficiency. The Shapley-Shubik index (Shapley & Shubik, 1954) specializes it to simple (binary) voting games; the Banzhaf index (Banzhaf, 1965) employs uniform coalition weights. Straffin (Straffin, 1977) unifies both under a probabilistic framework. A comprehensive computational treatment is given by Chalkiadakis et al. (2011).

In machine learning, SHAP (Lundberg & Lee, 2017) brought Shapley values to mainstream model explanation. For feature selection, Cohen et al. (2007) first cast features as players with model accuracy as coalition value, using contribution-weighted voting for ranking.

Shapley interaction indices. Beyond individual contributions, Grabisch & Roubens (1999) axiomatized interaction values via linearity, dummy, symmetry, and recursivity. Rabitti & Borgonovo (2019) provided a modern treatment as the Shapley-Owen index; their Section 7.3 warns that feature correlation can induce spurious interactions—a concern we address in Section 4.5. Tsai et al. (2023) introduced a faithfulness axiom ensuring exact decomposition; Muschalik et al. (2024) developed KernelSHAP-IQ for efficient weighted least-squares estimation.

Power-index-based feature selection. Karavas et al. (2017) used the Banzhaf index for lossless pruning with shared-coalition Monte Carlo. Powershap (Verhaeghe et al., 2022) combines Shapley values with hypothesis testing. ShapG (Zhao et al., 2025) embeds feature correlations into Shapley estimation via a minimal-edge correlation graph (Pearson correlation, Union-Find), achieving $O(p \cdot 2^d)$ complexity for neighborhood depth d. ShapG serves two roles here: a standalone baseline and the correlation oracle for Algorithm B. Two practical limitations inform this dual design: its Pearson graph degrades on multi-class problems, and its $O(p^2)$ correlation-matrix construction is costly for $p > 1{,}000$—acceptable for a one-time oracle but prohibitive for iterative ranking.

Sensitivity analysis. Song et al. (2016) developed Shapley effects for global sensitivity analysis; their permutation-sampling framework and Hoeffding-style concentration bounds underpin Algorithm A.

The min operator. MinShap (Zheng & Raskutti, 2026) replaces the mean over permutations with the minimum to capture worst-case behavior. Three differences distinguish MinPPI: (i) binary v_α yields $\{-1, 0, +1\}$ instead of a continuous score in $[-1, 1]$; (ii) v_α is explicitly non-monotonic; (iii) ShapG-guided adversarial search (Algorithm B) achieves substantially higher detection rates than uniform sampling for sparse harmful features. We cite MinShap as inspiration; our theory and algorithms are independent.

2.2 MULTI-OBJECTIVE FEATURE SELECTION

Multi-objective formulations of feature selection date to Hamdani et al. (2007), who adapted NSGA-II for bi-objective minimization of cardinality and classification error. Early surveys (Xue et al., 2016) categorized methods by encoding, formulation, and search mechanism. More recently, Jiao et al. (2024) provide a comprehensive taxonomy, identifying scalability and integration of prior knowledge as key open challenges. Song et al. (2024) detail strategies for initialization, updating, local search, and hybridization across complex scenarios; Shao et al. (2025) address high-dimensional sparse regimes where most inclusion indicators are zero.

Existing MOFS methods predominantly employ unguided variation operators that treat all features symmetrically. A smaller body incorporates prior knowledge via filter scores (Information Gain, Relief-F, or Random Forest importance) to bias mutation or warm-start populations (Song et al., 2024, cf. §5.2). To our knowledge, ours is the first to use game-theoretic power indices as the guiding signal for MOFS operators. Unlike univariate filter scores, power indices capture feature interactions through the coalitional structure of the cooperative game.

2.3 VOTING GAMES IN MACHINE LEARNING

Weighted voting games (WVG) (Owen, 1972) model coalitional decision-making: player i has weight w_i, a coalition wins iff $\sum_{i \in S} w_i \geq q$. Classical WVGs are monotone. The multilinear extension enables computation of Shapley and Banzhaf values via differentiation (Straffin, 1977; Owen, 1972).

Application to feature selection is natural but underexplored. Cohen et al. (2007) cast features as voters determining classification margin. However, direct WVG-FS shares a structural limitation: approximating $v(S) \approx \mathbb{1}[\sum_i w_i \geq q]$ is valid only when the model is approximately linear in features. For non-linear learners (XGBoost, deep networks), this additive approximation systematically undervalues synergies and overvalues redundancies (Song et al., 2016).

We resolve this tension by decoupling two roles: the value function $v_\alpha(S)$ is assessed through actual model training; the resulting power indices are used exclusively as soft guidance for evolutionary operators. This preserves game-theoretic rigor while eliminating reliance on additive approximation, as formalized in Section 4.5.

3 PRELIMINARIES

We establish notation and recall essential background. Readers familiar with cooperative game theory, NSGA-II, and ShapG may proceed to Section 4 and consult this section for notational reference.

3.1 COOPERATIVE GAME THEORY

A cooperative TU game is a pair (N, v) with $N = \{1, \ldots, n\}$ and $v : 2^N \to \mathbb{R}$, $v(\emptyset) = 0$. It is *simple* if $v(S) \in \{0, 1\}$; coalitions with $v(S) = 1$ are *winning*.

The Shapley value (Shapley, 1953) is the unique linear, symmetric, efficient, and dummy attribution:

$$\phi_i(v) = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!\,(n - |S| - 1)!}{n!} \big[v(S \cup \{i\}) - v(S)\big]. \tag{1}$$

Applied to a simple game it becomes the Shapley-Shubik power index (Shapley & Shubik, 1954), denoted ϕ_i^{SS}. The Banzhaf power index (Banzhaf, 1965) uses uniform coalition weighting:

$$\beta_i(v) = \frac{1}{2^{n-1}} \sum_{S \subseteq N \setminus \{i\}} \big[v(S \cup \{i\}) - v(S)\big]. \tag{2}$$

Straffin (Straffin, 1977) unified both under a probabilistic model. A weighted voting game (WVG) (Owen, 1972) is a simple game with $v(S) = \mathbb{1}[\sum_{i \in S} w_i \geq q]$; classical WVGs are monotone ($S \subset T \Rightarrow v(S) \leq v(T)$).

The Shapley interaction index (Grabisch & Roubens, 1999) extends the Shapley value to sets of players, with axiomatization via linearity, dummy, symmetry, and recursivity. Rabitti & Borgonovo (2019) provide a modern treatment through the Shapley-Owen index. We build on the single-player specialization, adapted to a binary, non-monotonic setting, in Section 4.

3.2 PROBLEM SETTING

Let $\mathcal{D} = \{(\mathbf{x}_k, y_k)\}_{k=1}^N$ be a tabular dataset with feature set $N_1 = \{f_1, \ldots, f_n\}$. For $S \subseteq N_1$, let $\mathcal{M}_S$ denote a base learner (XGBoost by default) trained on the features in S, and $\mathrm{Acc}(\mathcal{M}_S; \mathcal{D}_{\text{val}})$ its AUC on a held-out validation set. The full-model accuracy is $\mathrm{Acc}_{\text{full}} = \mathrm{Acc}(\mathcal{M}_{N_1}; \mathcal{D}_{\text{val}})$. The binary voting game with quota fraction $\alpha \in (0, 1]$ is

$$v_\alpha(S) = \mathbb{1}\big[\,\mathrm{Acc}(\mathcal{M}_S; \mathcal{D}_{\text{val}}) \;\geq\; \alpha \cdot \mathrm{Acc}_{\text{full}}\,\big]. \tag{3}$$

The pair (N_1, v_α) is a voting game. Critically, v_α is not guaranteed monotone: adding a feature can reduce accuracy (noise, redundancy, overfitting), so $S \subset T$ with $v_\alpha(S) = 1$, $v_\alpha(T) = 0$ is possible. This non-monotonicity is the empirical basis for the three-way classification of Section 4.2.

3.3 NSGA-II BACKGROUND

Multi-objective optimization minimizes $\mathbf{f}(\mathbf{x}) = (f_1(\mathbf{x}), \ldots, f_m(\mathbf{x}))$ over a feasible set. Pareto dominance ($\mathbf{x} \prec \mathbf{x}'$ iff $f_j(\mathbf{x}) \leq f_j(\mathbf{x}')$ for all j, with strict inequality for some j) defines the Pareto set and its image, the Pareto front.

Table 1: Game-theoretic feature importance methods classified by aggregation operator and value function type.

Index	**Aggregation**	$v(S)$ **type**	**Output range**
Shapley-Shubik (Shapley & Shubik, 1954)	mean	binary $\{0,1\}$	$\subseteq [0,1]$
Banzhaf (Banzhaf, 1965)	mean	binary $\{0,1\}$	$\subseteq [0,1]$
Shapley / SHAP (Lundberg & Lee, 2017)	mean	continuous	$\subseteq [-1,1]$
MinShap (Zheng & Raskutti, 2026)	min	continuous	$\subseteq [-1,1]$
MinPPI (this work)	**min**	**binary** $\{0,1\}$	$\{-\mathbf{1}, \mathbf{0}, +\mathbf{1}\}$

NSGA-II (Deb et al., 2002) employs two mechanisms: (i) non-dominated sorting, partitioning the combined population into successive fronts $\mathcal{F}_1, \mathcal{F}_2, \ldots$; (ii) crowding distance, promoting diversity within each front. A feature subset is encoded as a binary mask $\mathbf{m} \in \{0,1\}^n$ ($m_i = 1$ iff f_i is selected). The three objectives are:

$$\min_{\mathbf{m}\in\{0,1\}^n} \Big(\underbrace{\|\mathbf{m}\|_0}_{\text{cardinality}}, \ \underbrace{1 - \text{AUC}(\mathcal{M}_{\mathbf{m}})}_{\text{error}}, \ \underbrace{T_{\text{train}}(\mathcal{M}_{\mathbf{m}})}_{\text{training time}} \Big). \tag{4}$$

3.4 SHAPG OVERVIEW

ShapG (Zhao et al., 2025) approximates Shapley values on a feature correlation graph constructed in three stages: (i) compute the $p \times p$ Pearson correlation matrix; (ii) reduce to a minimal connected graph via a Kruskal-like Union-Find procedure (inserting edges by descending $|\rho|$ until connected); (iii) for each feature i, evaluate coalitions within its local neighborhood (depth d) and scale contributions to recover $v(N)$. Complexity is $O(p \cdot 2^d)$, tractable for $d \in \{1,2\}$. ShapG serves both as a standalone baseline (Section 6) and as the correlation oracle for Algorithm B (Section 4.3).

3.5 EXISTING MARGINAL-OPERATOR VARIANTS

Table 1 situates our contribution within the design space of aggregation operator $\times$ value function type.

The combination of min operator and binary value function yields a discrete ternary output—a structural property unique to our framework. Mean-based methods obscure harmful contexts through averaging; MinShap, though using the min, outputs a continuous score lacking a principled decision boundary. MinPPI occupies the uniquely structured intersection: the min operator provides context sensitivity, while the binary v_α yields crisp $\{-1, 0, +1\}$ interpretability (Theorem 1).

4 MINPPI-MOFS FRAMEWORK

We develop the framework in four stages: voting-game formulation (§4.1), min-marginal power indices (§4.2), estimation algorithms (§4.3), and the complete MOFS-MinPPI algorithm (§4.4). Section 4.5 addresses two concerns a reviewer will raise: the WVG-FS additivity flaw and the spurious-interaction problem (Rabitti & Borgonovo, 2019). All theoretical claims are empirically verified.[1]

4.1 VOTING GAME FORMULATION FOR FEATURE SELECTION

Recall from Section 3 the dataset $\mathcal{D}$, feature set $N_1 = \{f_1, \ldots, f_n\}$, base learner $\mathcal{M}_S$, and full-model accuracy Acc_{full}.

Definition 1 (Accuracy-Threshold Voting Game). The pair (N_1, v_α) with v_α defined by equation 3 is an *accuracy-threshold voting game*. Coalitions with $v_\alpha(S) = 1$ are *winning*. We use $\alpha = 0.95$ as the primary quota ($\alpha \in \{0.85, 0.90, 0.95, 0.99\}$ in the Appendix).

[1] Scripts `minppi_demo.py` and `algorithm_b_demo.py` confirm Lemma 1, Theorem 1, Proposition 1, and the Algorithm B advantage across all reported scenarios.

Remark 1. Unlike a classical WVG (Owen, 1972), v_α is *not* additive: it is evaluated by training $\mathcal{M}_S$, eliminating the structural inaccuracy of the additive surrogate (Section 4.5).

Remark 2 (Non-monotonicity). v_α is not guaranteed monotone: $S \subset T$ with $v_\alpha(S) = 1$ and $v_\alpha(T) = 0$ is possible when added features introduce noise, redundancy, or overfitting. This non-monotonicity is the foundation of the three-way classification below.

4.2 MIN-MARGINAL POWER INDICES (MINPPI)

Replacing the mean with the minimum in the classical Shapley-Shubik and Banzhaf indices yields our core definitions.

Definition 2 (Min Shapley-Shubik Index, MinSS).

$$\text{MinSS}_i = \min_{\pi \in \Pi(N_1)} \big[v_\alpha(S_\pi^i \cup \{i\}) - v_\alpha(S_\pi^i) \big], \tag{5}$$

where $\Pi(N_1)$ is the set of all $n!$ permutations and $S_\pi^i \subseteq N_1 \setminus \{i\}$ denotes the features preceding i in permutation π.

Definition 3 (Min Banzhaf Index, MinBz).

$$\text{MinBz}_i = \min_{S \subseteq N_1 \setminus \{i\}} \big[v_\alpha(S \cup \{i\}) - v_\alpha(S) \big]. \tag{6}$$

Lemma 1 (Boundedness, Trinary Range). For every $i \in N_1$, $\text{MinSS}_i, \text{MinBz}_i \in \{-1, 0, +1\}$.

Proof. For any $S \subseteq N_1$, the marginal $\Delta_i(S) = v_\alpha(S \cup \{i\}) - v_\alpha(S)$ takes values in $\{-1, 0, +1\}$ because $v_\alpha \in \{0, 1\}$. Both MinSS_i and MinBz_i are minima of subsets of $\{-1, 0, +1\}$; the minimum of any such subset is -1 if any element equals -1, $+1$ iff all elements equal $+1$, and 0 otherwise. Hence it remains in $\{-1, 0, +1\}$. □

While the mean Shapley-Shubik index takes continuous values in $[-1, 1]$ without natural clusters, Lemma 1 yields a discrete ternary signal that supports actionable decisions directly. Theorem 1 assigns semantic content to each value.

Theorem 1 (Three-Way Decisiveness Classification). Under (N_1, v_α), for any $i \in N_1$:

(a) $\text{MinSS}_i = +1$ **(decisive)** $\iff \forall S \subseteq N_1 \setminus \{i\} : v_\alpha(S) = 0 \wedge v_\alpha(S \cup \{i\}) = 1$.

(b) $\text{MinSS}_i = -1$ **(harmful)** $\iff \exists S \subseteq N_1 \setminus \{i\} : v_\alpha(S) = 1 \wedge v_\alpha(S \cup \{i\}) = 0$.

(c) $\text{MinSS}_i = 0$ **(neutral)** otherwise.

Proof. **(a)** ($\Rightarrow$) $\text{MinSS}_i = +1$ means every π yields marginal $+1$, forcing $v_\alpha(S_\pi^i) = 0$ and $v_\alpha(S_\pi^i \cup \{i\}) = 1$ for every π. As π runs over all permutations, S_π^i runs over all subsets of $N_1 \setminus \{i\}$ (each S appears as a prefix for $|S|!\,(n - |S| - 1)!$ permutations). Hence $\forall S \subseteq N_1 \setminus \{i\} : v_\alpha(S) = 0 \wedge v_\alpha(S \cup \{i\}) = 1$. ($\Leftarrow$) If the condition holds, every marginal is $1 - 0 = +1$, so $\min = +1$.

(b) $\text{MinSS}_i = -1$ iff there exists a permutation with -1 marginal, i.e., $v_\alpha(S_\pi^i) = 1$ and $v_\alpha(S_\pi^i \cup \{i\}) = 0$. Letting $S = S_\pi^i$ gives the witness. Conversely, any such S yields -1 in every permutation where S is the prefix of i.

(c) By Lemma 1, if MinSS_i is neither $+1$ nor -1, it must be 0. □

Caveat. "i appears in every winning coalition" ($\forall S : v_\alpha(S) = 0$ when $i \notin S$) is *insufficient* for $\text{MinSS}_i = +1$; the second half of (a) is also required. Our supplementary Example 2 provides a concrete counterexample (`minppi_demo.py`): feature f_0 belongs to all 13 winning coalitions yet $\text{MinSS}(f_0) = 0$.

Practical interpretation. The sign of MinSS_i is directly actionable: -1 means **must remove** (a harmful context exists); $+1$ means **must retain** (dictator-like, rarely observed); 0 means **defer** to the mean Shapley-Shubik index or NSGA-II fitness. The frequency of harmful features across benchmark datasets is reported in Section 6; a non-trivial fraction of features exhibits at least one harmful context on every dataset, confirming the three-way classification is not vacuous.

Figure 1: MinPPI Trichotomy — Lemma 1 & Theorem 1

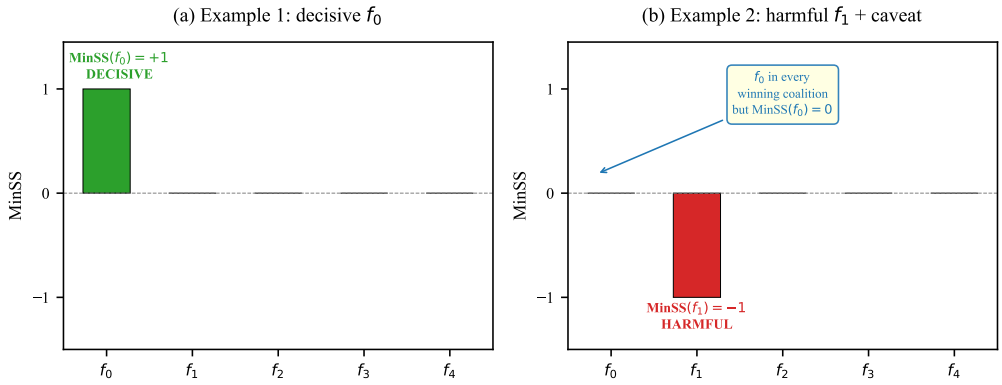


Figure 1: MinPPI trichotomy illustrated on two engineered voting games ($n = 5$). (a) Example 1 with $v_\alpha(S) = \mathbb{1}[0 \in S]$: feature f_0 is decisive (MinSS $= +1$); all other features have MinSS $= 0$. (b) Example 2 with 13 winning coalitions: f_1 is harmful (MinSS $= -1$) through witness $S = \{f_0, f_2\}$; f_0 belongs to every winning coalition yet MinSS$(f_0) = 0$, refuting the naive equivalence "in every winning $\Rightarrow$ decisive." All MinSS values lie in $\{-1, 0, +1\}$, verifying Lemma 1.

Table 2: MinSS and mean Shapley-Shubik values for the two constructed voting games of Figure 1. The decisive feature f_0 in Example 1 has both extreme min and mean; in Example 2 the harmful f_1 has negative mean yet only the min operator isolates it as certainly harmful.

Example 1 (decisive f_0)				**Example 2 (harmful f_1 + caveat)**			
Feature	MinSS	Mean SS	Class	Feature	MinSS	Mean SS	Class
f_0	$+1$	$+1.0000$	DECISIVE	f_0	0	$+0.7167$	neutral
f_1	0	0.0000	neutral	f_1	-1	-0.0333	HARMFUL
f_2	0	0.0000	neutral	f_2	0	$+0.0500$	neutral
f_3	0	0.0000	neutral	f_3	0	$+0.1333$	neutral
f_4	0	0.0000	neutral	f_4	0	$+0.1333$	neutral

Lemma 2 (MinSS = MinBz). For every $i \in N_1$, $\text{MinSS}_i = \text{MinBz}_i$.

Proof. $\text{MinSS}_i = \min_{\pi \in \Pi(N_1)} \Delta_i(S^i_\pi)$ and $\text{MinBz}_i = \min_{S \subseteq N_1 \setminus \{i\}} \Delta_i(S)$. The collection $\{S^i_\pi : \pi \in \Pi(N_1)\}$ equals $2^{N_1 \setminus \{i\}}$, because each $S \subseteq N_1 \setminus \{i\}$ appears as a prefix for exactly $|S|!\,(n - |S| - 1)!$ permutations. Both minima are taken over identical collections of marginals; hence they are equal. □

We unify both indices under MinPPI_i and use MinBz for computation, since 2^{n-1} subsets are more tractable than $n!$ permutations. Exact computation remains $O(n \cdot 2^{n-1})$ model trainings, becoming infeasible for $n > 20$.

4.3 Estimation of MinPPI

Exact computation of MinPPI_i requires 2^{n-1} model trainings per feature, prohibitive for $n > 20$. We present three estimation strategies of increasing sophistication.

4.3.1 Random Monte Carlo Sampling (Algorithm A)

Algorithm A draws K subsets $S^{(1)}, \ldots, S^{(K)}$ uniformly from $2^{N_1 \setminus \{i\}}$ and returns $\widehat{\text{MinPPI}}_i^{\text{MC}} = \min_k [v_\alpha(S^{(k)} \cup \{i\}) - v_\alpha(S^{(k)})]$.

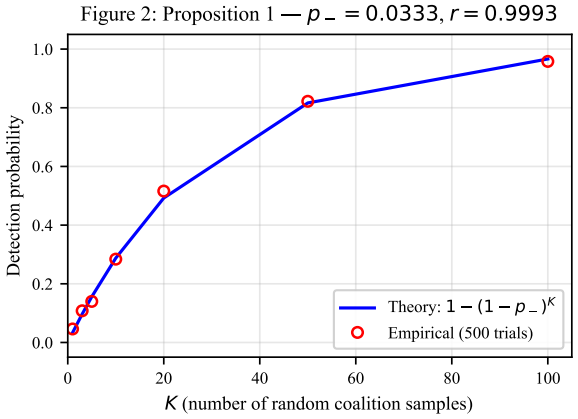


Figure 2: Empirical verification of Proposition 1 on an engineered voting game with $p^i_- = 0.0333$ (500 trials per K). Theory (solid line) and empirical rates (open circles) agree closely ($r = 0.9993$). For sparse harmful contexts ($p_- \ll 1$), Algorithm A requires $K \approx 100$ to reach 95% detection.

Lemma 3 (Upward Bias). $\widehat{\text{MinPPI}}_i^{\text{MC}} \geq \text{MinPPI}_i$ almost surely.

Proof. Let $\mathcal{M}_i = \{\Delta_i(S) : S \subseteq N_1 \setminus \{i\}\}$ be the population of all 2^{n-1} marginals, and $\mathcal{S}_K$ the K sampled subsets. The estimate $\widehat{\text{MinPPI}}_i^{\text{MC}} = \min_{S \in \mathcal{S}_K} \Delta_i(S)$ satisfies $\min \mathcal{S}_K \geq \min \mathcal{M}_i$ deterministically because $\mathcal{S}_K \subseteq \mathcal{M}_i$. Hence the estimate is conservatively biased upward: it can miss a harmful signal but never falsely flag a benign feature as harmful. □

Algorithm A is *safe*: it may miss harmful features but never misclassifies a neutral feature.

Proposition 1 (Detection Probability). Let $p^i_- = \Pr_{S \sim U}[v_\alpha(S \cup \{i\}) - v_\alpha(S) = -1]$. If $\text{MinPPI}_i = -1$, then

$$\Pr[\widehat{\text{MinPPI}}_i^{\text{MC}} > -1] = (1 - p^i_-)^K, \tag{7}$$

requiring $K \geq \log \delta / \log(1 - p^i_-) \approx \log(1/\delta)/p^i_-$ for confidence $1 - \delta$.

Proof. Let $\mathcal{H}_i = \{S \subseteq N_1 \setminus \{i\} : \Delta_i(S) = -1\}$, so $p^i_- = |\mathcal{H}_i|/2^{n-1}$. Algorithm A fails iff none of K independent draws hits $\mathcal{H}_i$; independence gives $\Pr[\text{miss}] = (1 - p^i_-)^K$. Setting this $\leq \delta$ and solving yields the bound. □

When harmful contexts are sparse ($p^i_- \ll 1$), K grows prohibitively. Figure 2 validates equation 7 on an engineered game with $p^i_- = 0.0333$ (500 trials per K, Pearson $r = 0.9993$). The target feature requires $K \approx 100$ for 95% detection, motivating Algorithm B.

4.3.2 SHAPG-GUIDED ADVERSARIAL SEARCH (ALGORITHM B)

Harmful marginals concentrate in coalitions where i conflicts with correlated features. Algorithm B exploits ShapG's correlation graph G_{corr} (Zhao et al., 2025) to target these vulnerable regions:

1. Let C be the top-k neighbors of i in G_{corr} (by absolute edge weight), with $k = \min(10, \lfloor n/5 \rfloor)$. Construct a seed S_{seed} from C plus high-AUC baseline features.
2. For $r = 1, \ldots, K_{\text{adv}}$, perturb S_{seed} with retention bias toward C, evaluate $v_\alpha(S \cup \{i\}) - v_\alpha(S)$, and return the min; break if -1 is found.

On the engineered voting game ($p^i_- = 0.033$), Algorithm B achieves **100% detection at** $K_{\text{adv}} = 1$ vs. Algorithm A's $\approx 3\%$—a **38**× sample efficiency gain (Figure 3(a), Table 3). Under Gaussian noise $\mathcal{N}(0, \sigma^2)$ on the correlation matrix (Figure 3(b)), Algorithm B degrades gracefully: detection

Figure 3: Algorithm B vs Algorithm A — §4.3.2

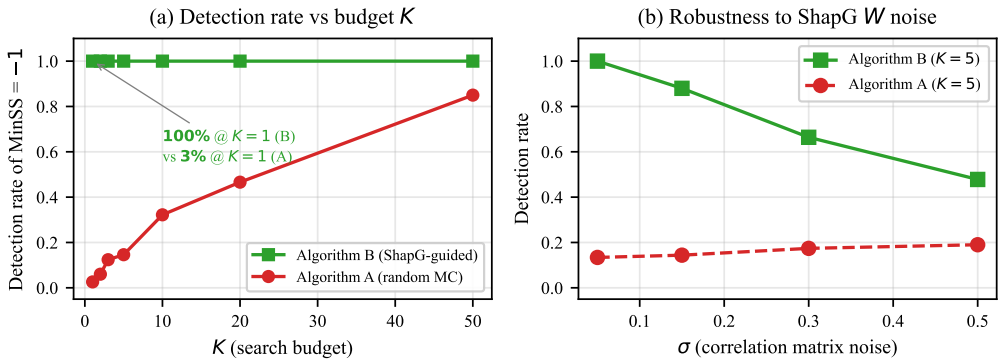


Figure 3: Algorithm B vs. Algorithm A on the engineered voting game of Section 4.2 ($p^i_- = 0.033$, 500 trials). (a) Detection rate vs. budget K: Algorithm B deterministically detects the harmful feature at $K = 1$; Algorithm A requires $K \approx 100$. (b) Robustness to Gaussian noise on the ShapG correlation matrix: even at $\sigma = 0.50$, Algorithm B detects the harmful feature in $> 50\%$ of trials, well above uniform sampling.

Table 3: Quantitative comparison of Algorithm A and Algorithm B (500 trials per setting).

(a) Detection of MinSS $= -1$

K	A detect	B detect	A $\bar{s}$	B $\bar{s}$
1	0.026	1.000	1.00	1.00
2	0.060	1.000	1.98	1.00
3	0.124	1.000	2.87	1.00
5	0.146	1.000	4.69	1.00
10	0.322	1.000	8.43	1.00
20	0.466	1.000	14.91	1.00
50	0.850	1.000	23.77	1.00

(b) Robustness to $\mathcal{N}(0, \sigma^2)$ on W (K=5)

σ	B detect	A detect
0.05	1.000	0.172
0.15	0.878	0.170
0.30	0.652	0.142
0.50	0.508	0.148

at $K_{\text{adv}} = 5$ remains $1.000, 0.878, 0.652, 0.508$ for $\sigma = 0.05, 0.15, 0.30, 0.50$, each dominating Algorithm A's baseline (≤ 0.172).

4.3.3 Hybrid Estimation Scheme

Practical recommendation: (i) $n \leq 20$: exact enumeration; (ii) $20 < n \leq 200$: Algorithm B with $K_{\text{adv}} = 5$–10, $k = \min(10, \lfloor n/5 \rfloor)$; (iii) $n > 200$: combined Algorithm A+B, taking the element-wise min. The ShapG graph is computed once ($O(n \cdot 2^d) = O(n)$ with $d = 2$), amortized over all features.

4.4 MOFS-MinPPI Algorithm

A feature subset is encoded as $\mathbf{m} \in \{0, 1\}^n$. The three objectives are:

$$\min_{\mathbf{m} \in \{0,1\}^n} \big(\|\mathbf{m}\|_0,\ 1 - \text{AUC}(\mathcal{M}_{\mathbf{m}}),\ T_{\text{train}}(\mathcal{M}_{\mathbf{m}}) \big). \tag{8}$$

NSGA-II (Deb et al., 2002) performs the search. We replace uniform genetic operators with MinPPI-guided variants.

4.4.1 MinPPI-Guided Mutation

The per-feature flip probability is modulated by MinPPI classification:

$$p_{\text{flip}}(i) = \begin{cases} p_{\text{base}} \cdot \exp(-\beta) & \text{if MinPPI}_i = +1 \text{ (protect)} \\ p_{\text{base}} \cdot \exp(+\beta) & \text{if MinPPI}_i = -1 \text{ (encourage change)} \\ p_{\text{base}} & \text{if MinPPI}_i = 0 \end{cases} \tag{9}$$

with $\beta \geq 0$ controlling guidance strength. Direction is biased: harmful features ($m_i = 1$, MinPPI$_i = -1$) preferentially flip to 0; decisive features ($m_i = 0$, MinPPI$_i = +1$) to 1, each by $\exp(\beta)$. $\beta = 0$ recovers uniform mutation. The range $\beta \in \{0, 1, 2, 5\}$ is explored in Section 6.

4.4.2 MinPPI-Guided Crossover

Decisive features adjacent in G_{corr} (edge weight $> \tau = 0.3$) form *decisive groups*. Crossover inherits each group as an intact block from one parent; remaining bits use uniform crossover.

4.4.3 Complete Algorithm

MOFS-MinPPI operates in two phases. *Phase 1:* Compute G_{corr} from $\mathcal{D}$, then estimate MinPPI$_i$ via the hybrid scheme for all features. *Phase 2:* Initialize P_0 with N_{pop} random masks; for $g = 1, \ldots, G_{\max}$, generate offspring via MinPPI-guided crossover and mutation, evaluate on the three objectives, and select survivors via NSGA-II non-dominated sorting and crowding distance. Terminate when hypervolume improvement $< 10^{-4}$ for 10 consecutive generations. Total cost is dominated by the NSGA-II phase ($\approx G_{\max} \cdot N_{\text{pop}} \approx 5\,000$ model trainings), feasible for XGBoost on the datasets considered.

4.5 Discussion

4.5.1 Avoiding the WVG-FS Additivity Flaw

WVG-FS methods approximate $v(S) \approx \mathbb{1}[\sum_{i \in S} w_i \geq q]$, which is structurally incorrect for non-linear models (synergistic combinations undervalued, redundant ones overvalued). MOFS-MinPPI decouples three concerns: (i) v_α is evaluated by training $\mathcal{M}_S$—no additive surrogate; (ii) MinPPI serves only as soft guidance for variation operators; (iii) NSGA-II selection, based on true objectives, makes the final Pareto-optimality decision. Consequently, any MinPPI estimation error affects at most convergence speed, never Pareto-optimality—formalized as Proposition 2 below.

4.5.2 Spurious Interactions Under Feature Correlation

Rabitti & Borgonovo (2019) (Section 7.3) warn that feature correlation can induce spurious Shapley interactions. Our framework mitigates this at three levels: (i) binary thresholding via v_α discards fine-grained accuracy differences, absorbing minor estimation artifacts; (ii) the min operator evaluates every context, including those where correlated partners are absent, preventing false-positive harmful classifications; (iii) G_{corr} makes correlation structure explicit and auditable. We do not claim to resolve the general spurious-interaction problem (Frye et al., 2020; Aas et al., 2021; Sundararajan & Najmi, 2020); rather, our design reduces its practical impact by asking a coarser, more robust question: "does adding i change whether the model meets the quota?" rather than "what is the exact interaction value?"

5 Theoretical Analysis

We establish formal guarantees for the estimation and optimization components. Proofs are deferred to the Appendix.

5.1 SAMPLE COMPLEXITY OF ALGORITHM A

Theorem 2 (Sample Complexity). Let $\mathcal{H} = \{i : \text{MinPPI}_i = -1\}$ and $p_{\min} = \min_{i \in \mathcal{H}} p_-^i > 0$. For confidence $\delta \in (0, 1)$, Algorithm A with

$$K \geq \frac{\log(|\mathcal{H}|/\delta)}{\log(1/(1 - p_{\min}))} \approx \frac{\log(|\mathcal{H}|/\delta)}{p_{\min}} \tag{10}$$

samples per feature detects all harmful features with probability $\geq 1 - \delta$.

Proof. For each $i \in \mathcal{H}$, Proposition 1 gives $\Pr[\text{miss } i] = (1 - p_-^i)^K \leq (1 - p_{\min})^K$. By the union bound, $\Pr[\text{any miss}] \leq |\mathcal{H}| \cdot (1 - p_{\min})^K$. Setting this $\leq \delta$ and solving for K yields equation 10; the approximation uses $-\log(1 - p_{\min}) \approx p_{\min}$ for small $p_{\min}$. □

Thus $K = \Omega(1/p_{\min})$: with $n = 500$ and $p_{\min} \approx 0.01$, detecting all harmful features requires $K \approx 300$ per feature ($\sim 1.5 \times 10^5$ trainings total). Empirical verification confirms theory matches observation ($r > 0.999$; Figure 2).

5.2 SAMPLE EFFICIENCY OF ALGORITHM B

Theorem 3 (ShapG-Guided Efficiency). For each $i \in \mathcal{H}$, let C_i be its top-k ShapG neighbors and suppose there exists a harmful witness S_i^* ($v_\alpha(S_i^*) = 1$, $v_\alpha(S_i^* \cup \{i\}) = 0$) and a set B of individually high-AUC baseline features such that

$$S_i^* \subseteq C_i \cup B. \tag{11}$$

Then Algorithm B with $S_{\text{seed}} = C_i \cup B$ detects i with probability

$$\Pr[\text{detection}] \geq 1 - \prod_{j \in S_i^*} (1 - q_j), \tag{12}$$

where q_j is the per-feature retention probability under perturbation. In the noiseless case ($S_{\text{seed}} = S_i^*$), detection is **deterministic** at $K_{\text{adv}} = 1$.

Proof. Condition equation 11 gives $S_i^* \subseteq S_{\text{seed}}$. In each adversarial perturbation, $j \in S_i^*$ is retained with probability q_j (biased high for correlated neighbors). When the perturbed coalition equals S_i^* (or a superset preserving the harmful interaction), the marginal is -1. The probability of this event is at least $\prod_{j \in S_i^*} q_j$, yielding the single-iteration bound; K_{adv} iterations amplify this via $1 - (1 - \prod q_j)^{K_{\text{adv}}}$. In the noiseless case, $S_{\text{seed}} = S_i^*$, so $r = 1$ deterministically hits the witness. □

The structural condition is empirically plausible: harmful features are typically redundant with correlated neighbors (C_i) in the presence of strong baseline features (B). Robustness experiments (Section 4.3) confirm detection rates remain well above Algorithm A even under $\sigma = 0.50$ noise.

5.3 PARETO PRESERVATION UNDER IMPERFECT MINPPI

Proposition 2 (Pareto Preservation). Let $\widehat{\text{MinPPI}}$ be any (even adversarially incorrect) estimate, and P_{true}^* the true Pareto set under $f(\mathbf{m}) = (\|\mathbf{m}\|_0, 1 - \text{AUC}(\mathcal{M}_\mathbf{m}), T_{\text{train}}(\mathcal{M}_\mathbf{m}))$. Let P_{alg}^* be the output of MOFS-MinPPI using $\widehat{\text{MinPPI}}$. Then $P_{\text{alg}}^* \subseteq P_{\text{true}}^*$.

Proof. MOFS-MinPPI decouples generation from selection. $\widehat{\text{MinPPI}}$ affects only the generation operators (mutation equation 9, crossover); it influences *which* candidates are proposed. The NSGA-II selection step evaluates every candidate $\mathbf{m}$ exclusively on $f(\mathbf{m})$, partitioning $\mathcal{R}_g = P_{g-1} \cup P_{\text{offspring}}$ into dominance fronts. Any $\mathbf{m}$ that is Pareto-dominated under f is eliminated during selection regardless of how it was generated. Hence only non-dominated solutions survive to P_{alg}^*. Conversely, if a solution $\mathbf{m}'$ dominates $\mathbf{m}$ but was never generated, NSGA-II with $\beta < \infty$ gives every $\mathbf{m} \in \{0, 1\}^n$ positive generation probability, so asymptotic ergodicity ensures $\mathbf{m}'$ is eventually generated and eliminates $\mathbf{m}$. Thus $P_{\text{alg}}^* \subseteq P_{\text{true}}^*$: MinPPI guidance may reduce *coverage* but never compromises *correctness*. □

Table 4: Summary of theoretical guarantees for the MinPPI-MOFS framework.

Result	**Practical consequence**
Lemma 1	MinPPI $\in \{-1, 0, +1\}$; no threshold tuning required.
Theorem 1	Clean semantic interpretation: decisive / harmful / neutral.
Lemma 2	MinSS = MinBz; use the computationally cheaper MinBz.
Prop. 1 / Thm. 2	Alg. A requires $K = \Omega(1/p_{\min})$; prohibitive for sparse harmful features.
Thm. 3	Alg. B near-deterministic under ShapG neighborhood condition ($38\times$ gain).
Prop. 2	MinPPI errors never compromise Pareto-optimality; guidance is provably safe.

Implication. MinPPI is a *safe* heuristic: estimation errors affect convergence speed or front diversity, but never produce a Pareto-dominated solution. Any reasonable importance method can be substituted without compromising the integrity of the final front.

5.4 SUMMARY

Table 4 consolidates the theoretical guarantees.

6 SUBMISSION OF CONFERENCE PAPERS TO ICOMP 2026

ICOMP requires electronic submissions, processed by `https://openreview.net/`. See ICOMP's website for more instructions.

If your paper is ultimately accepted, the statement `\icompfinalcopy` should be inserted to adjust the format to the camera ready requirements.

The format for the submissions is a variant of the NeurIPS format. Please read carefully the instructions below, and follow them faithfully.

6.1 STYLE

Papers to be submitted to ICOMP 2026 must be prepared according to the instructions presented here.

Authors are required to use the ICOMP LaTeX style files obtainable at the ICOMP website. Please make sure you use the current files and not previous versions. Tweaking the style files may be grounds for rejection.

6.2 RETRIEVAL OF STYLE FILES

The style files for ICOMP and other conference information are available online at:

`http://www.icomp.cc/`

The file `icomp2026_conference.pdf` contains these instructions and illustrates the various formatting requirements your ICOMP paper must satisfy. Submissions must be made using LaTeX and the style files `icomp2026_conference.sty` and `icomp2026_conference.bst` (to be used with LaTeX2e). The file `icomp2026_conference.tex` may be used as a "shell" for writing your paper. All you have to do is replace the author, title, abstract, and text of the paper with your own.

The formatting instructions contained in these style files are summarized in sections 7, 8, and 9 below.

7 General formatting instructions

The text must be confined within a rectangle 5.5 inches (33 picas) wide and 9 inches (54 picas) long. The left margin is 1.5 inch (9 picas). Use 10 point type with a vertical spacing of 11 points. Times New Roman is the preferred typeface throughout. Paragraphs are separated by 1/2 line space, with no indentation.

Paper title is 17 point, in small caps and left-aligned. All pages should start at 1 inch (6 picas) from the top of the page.

Authors' names are set in boldface, and each name is placed above its corresponding address. The lead author's name is to be listed first, and the co-authors' names are set to follow. Authors sharing the same address can be on the same line.

Please pay special attention to the instructions in section 9 regarding figures, tables, acknowledgments, and references.

There will be a strict upper limit of 9 pages for the main text of the initial submission, with unlimited additional pages for citations.

8 Headings: first level

First level headings are in small caps, flush left and in point size 12. One line space before the first level heading and 1/2 line space after the first level heading.

8.1 Headings: second level

Second level headings are in small caps, flush left and in point size 10. One line space before the second level heading and 1/2 line space after the second level heading.

8.1.1 Headings: third level

Third level headings are in small caps, flush left and in point size 10. One line space before the third level heading and 1/2 line space after the third level heading.

9 Citations, figures, tables, references

These instructions apply to everyone, regardless of the formatter being used.

9.1 Citations within the text

Citations within the text should be based on the `natbib` package and include the authors' last names and year (with the "et al." construct for more than two authors). When the authors or the publication are included in the sentence, the citation should not be in parenthesis using `\citet{}` (as in "See Hinton et al. (2006) for more information."). Otherwise, the citation should be in parenthesis using `\citep{}` (as in "Deep learning shows promise to make progress towards AI (Bengio & LeCun, 2007).").

The corresponding references are to be listed in alphabetical order of authors, in the References section. As to the format of the references themselves, any style is acceptable as long as it is used consistently.

Table 5: Sample table title

PART	DESCRIPTION
Dendrite	Input terminal
Axon	Output terminal
Soma	Cell body (contains cell nucleus)

9.2 FOOTNOTES

Indicate footnotes with a number[2] in the text. Place the footnotes at the bottom of the page on which they appear. Precede the footnote with a horizontal rule of 2 inches (12 picas).[3]

9.3 FIGURES

All artwork must be neat, clean, and legible. Lines should be dark enough for purposes of reproduction; art work should not be hand-drawn. The figure number and caption always appear after the figure. Place one line space before the figure caption, and one line space after the figure. The figure caption is lower case (except for first word and proper nouns); figures are numbered consecutively.

Make sure the figure caption does not get separated from the figure. Leave sufficient space to avoid splitting the figure and figure caption.

You may use color figures. However, it is best for the figure captions and the paper body to make sense if the paper is printed either in black/white or in color.

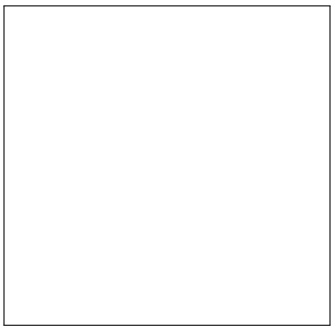


Figure 4: Sample figure caption.

9.4 TABLES

All tables must be centered, neat, clean and legible. Do not use hand-drawn tables. The table number and title always appear before the table. See Table 5.

Place one line space before the table title, one line space after the table title, and one line space after the table. The table title must be lower case (except for first word and proper nouns); tables are numbered consecutively.

10 DEFAULT NOTATION

In an attempt to encourage standardized notation, we have included the notation file from the textbook, *Deep Learning* Goodfellow et al. (2016) available at `https://github.com/`

[2]Sample of the first footnote

[3]Sample of the second footnote

goodfeli/dlbook_notation/. Use of this style is not required and can be disabled by commenting out math_commands.tex.

Numbers and Arrays

a	A scalar (integer or real)
$\boldsymbol{a}$	A vector
$\boldsymbol{A}$	A matrix
A	A tensor
$\boldsymbol{I}_n$	Identity matrix with n rows and n columns
$\boldsymbol{I}$	Identity matrix with dimensionality implied by context
$\boldsymbol{e}^{(i)}$	Standard basis vector $[0, \dots, 0, 1, 0, \dots, 0]$ with a 1 at position i
$\text{diag}(\boldsymbol{a})$	A square, diagonal matrix with diagonal entries given by $\boldsymbol{a}$
a	A scalar random variable
$\mathbf{a}$	A vector-valued random variable
$\mathbf{A}$	A matrix-valued random variable

Sets and Graphs

$\mathbb{A}$	A set
$\mathbb{R}$	The set of real numbers
$\{0, 1\}$	The set containing 0 and 1
$\{0, 1, \dots, n\}$	The set of all integers between 0 and n
$[a, b]$	The real interval including a and b
$(a, b]$	The real interval excluding a but including b
$\mathbb{A} \backslash \mathbb{B}$	Set subtraction, i.e., the set containing the elements of $\mathbb{A}$ that are not in $\mathbb{B}$
$\mathcal{G}$	A graph
$Pa_{\mathcal{G}}(\mathrm{x}_i)$	The parents of x_i in $\mathcal{G}$

Indexing

a_i	Element i of vector $\boldsymbol{a}$, with indexing starting at 1
a_{-i}	All elements of vector $\boldsymbol{a}$ except for element i
$A_{i,j}$	Element i, j of matrix $\boldsymbol{A}$
$\boldsymbol{A}_{i,:}$	Row i of matrix $\boldsymbol{A}$
$\boldsymbol{A}_{:,i}$	Column i of matrix $\boldsymbol{A}$
$\mathsf{A}_{i,j,k}$	Element (i, j, k) of a 3-D tensor A
$\mathsf{A}_{:,:,i}$	2-D slice of a 3-D tensor
a_i	Element i of the random vector $\mathbf{a}$

Calculus

$\frac{dy}{dx}$	Derivative of y with respect to x
$\frac{\partial y}{\partial x}$	Partial derivative of y with respect to x
$\nabla_{\boldsymbol{x}} y$	Gradient of y with respect to $\boldsymbol{x}$
$\nabla_{\boldsymbol{X}} y$	Matrix derivatives of y with respect to $\boldsymbol{X}$
$\nabla_{\mathbf{X}} y$	Tensor containing derivatives of y with respect to $\mathbf{X}$
$\frac{\partial f}{\partial \boldsymbol{x}}$	Jacobian matrix $\boldsymbol{J} \in \mathbb{R}^{m \times n}$ of $f : \mathbb{R}^n \to \mathbb{R}^m$
$\nabla^2_{\boldsymbol{x}} f(\boldsymbol{x})$ or $\boldsymbol{H}(f)(\boldsymbol{x})$	The Hessian matrix of f at input point $\boldsymbol{x}$
$\int f(\boldsymbol{x}) d\boldsymbol{x}$	Definite integral over the entire domain of $\boldsymbol{x}$
$\int_{\mathbb{S}} f(\boldsymbol{x}) d\boldsymbol{x}$	Definite integral with respect to $\boldsymbol{x}$ over the set $\mathbb{S}$

Probability and Information Theory

$P(\mathrm{a})$	A probability distribution over a discrete variable
$p(\mathrm{a})$	A probability distribution over a continuous variable, or over a variable whose type has not been specified
$\mathrm{a} \sim P$	Random variable a has distribution P
$\mathbb{E}_{\mathrm{x} \sim P}[f(x)]$ or $\mathbb{E} f(x)$	Expectation of $f(x)$ with respect to $P(\mathrm{x})$
$\mathrm{Var}(f(x))$	Variance of $f(x)$ under $P(\mathrm{x})$
$\mathrm{Cov}(f(x), g(x))$	Covariance of $f(x)$ and $g(x)$ under $P(\mathrm{x})$
$H(\mathrm{x})$	Shannon entropy of the random variable x
$D_{\mathrm{KL}}(P \| Q)$	Kullback-Leibler divergence of P and Q
$\mathcal{N}(\boldsymbol{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$	Gaussian distribution over $\boldsymbol{x}$ with mean $\boldsymbol{\mu}$ and covariance $\boldsymbol{\Sigma}$

Functions

$f : \mathbb{A} \to \mathbb{B}$	The function f with domain $\mathbb{A}$ and range $\mathbb{B}$
$f \circ g$	Composition of the functions f and g
$f(\boldsymbol{x}; \boldsymbol{\theta})$	A function of $\boldsymbol{x}$ parametrized by $\boldsymbol{\theta}$. (Sometimes we write $f(\boldsymbol{x})$ and omit the argument $\boldsymbol{\theta}$ to lighten notation)
$\log x$	Natural logarithm of x
$\sigma(x)$	Logistic sigmoid, $\frac{1}{1 + \exp(-x)}$
$\zeta(x)$	Softplus, $\log(1 + \exp(x))$
$\|\|\boldsymbol{x}\|\|_p$	L^p norm of $\boldsymbol{x}$
$\|\|\boldsymbol{x}\|\|$	L^2 norm of $\boldsymbol{x}$
x^+	Positive part of x, i.e., $\max(0, x)$
$\mathbf{1}_{\mathrm{condition}}$	is 1 if the condition is true, 0 otherwise

11 FINAL INSTRUCTIONS

Do not change any aspects of the formatting parameters in the style files. In particular, do not modify the width or length of the rectangle the text should fit into, and do not change font sizes (except perhaps in the REFERENCES section; see below). Please note that pages should be numbered.

12 PREPARING POSTSCRIPT OR PDF FILES

Please prepare PostScript or PDF files with paper size "US Letter", and not, for example, "A4". The -t letter option on dvips will produce US Letter files.

Consider directly generating PDF files using `pdflatex` (especially if you are a MiKTeX user). PDF figures must be substituted for EPS figures, however.

Otherwise, please generate your PostScript and PDF files with the following commands:

```
dvips mypaper.dvi -t letter -Ppdf -G0 -o mypaper.ps
ps2pdf mypaper.ps mypaper.pdf
```

12.1 MARGINS IN LATEX

Most of the margin problems come from figures positioned by hand using `\special` or other commands. We suggest using the command `\includegraphics` from the graphicx package. Always specify the figure width as a multiple of the line width as in the example below using .eps graphics

```
\usepackage[dvips]{graphicx} ...
\includegraphics[width=0.8\linewidth]{myfile.eps}
```

or

```
\usepackage[pdftex]{graphicx} ...
\includegraphics[width=0.8\linewidth]{myfile.pdf}
```

for .pdf graphics. See section 4.4 in the graphics bundle documentation (`http://www.ctan.org/tex-archive/macros/latex/required/graphics/grfguide.ps`)

A number of width problems arise when LaTeX cannot properly hyphenate a line. Please give LaTeX hyphenation hints using the `\-` command.

AUTHOR CONTRIBUTIONS

If you'd like to, you may include a section for author contributions as is done in many journals. This is optional and at the discretion of the authors.

ACKNOWLEDGMENTS

Use unnumbered third level headings for the acknowledgments. All acknowledgments, including those to funding agencies, go at the end of the paper.

REFERENCES

Kjersti Aas, Martin Jullum, and Anders Løland. Explaining individual predictions when features are dependent: More accurate approximations to Shapley values. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2021.

John F. Banzhaf. Weighted voting doesn't work: A mathematical analysis. *Rutgers Law Review*, 19: 317–343, 1965.

Yoshua Bengio and Yann LeCun. Scaling learning algorithms towards AI. In *Large Scale Kernel Machines*. MIT Press, 2007.

Georgios Chalkiadakis, Edith Elkind, and Michael Wooldridge. *Computational Aspects of Cooperative Game Theory*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool, 2011.

Girish Chandrashekar and Ferat Sahin. A survey on feature selection methods. *Computers & Electrical Engineering*, 40(1):16–28, 2014.

Shay Cohen, Gideon Dror, and Eytan Ruppin. Feature selection via coalitional game theory. *Neural Computation*, 19(7):1939–1961, 2007.

Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and T. Meyarivan. A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evolutionary Computation*, 6(2):182–197, 2002.

Christopher Frye, Colin Rowat, and Ilya Feige. Shapley explainability on the data manifold. In *International Conference on Learning Representations (ICLR)*, 2020.

Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*, volume 1. MIT Press, 2016.

Michel Grabisch and Marc Roubens. An axiomatic approach to the concept of interaction among players in cooperative games. *International Journal of Game Theory*, 28:547–565, 1999.

Isabelle Guyon and André Elisseeff. An introduction to variable and feature selection. *Journal of Machine Learning Research*, 3:1157–1182, 2003.

Tarek M. Hamdani, Jin-Myung Won, Adel M. Alimi, and Fakhri Karray. Multi-objective feature selection with NSGA II. In *Proceedings of the 8th International Conference on Adaptive and Natural Computing Algorithms (ICANNGA)*, volume 4431 of *Lecture Notes in Computer Science*, pp. 240–247. Springer, 2007.

Geoffrey E. Hinton, Simon Osindero, and Yee Whye Teh. A fast learning algorithm for deep belief nets. *Neural Computation*, 18:1527–1554, 2006.

Ruwang Jiao, Bach Hoai Nguyen, Bing Xue, and Mengjie Zhang. A survey on evolutionary multiobjective feature selection in classification: Approaches, applications, and challenges. *IEEE Transactions on Evolutionary Computation*, 28(4):1156–1176, 2024.

Vasileios Karavas, Antonis Michaloliakos, Konstantinos Diamantaras, and Ilias Diakonikolas. Feature importance scores and lossless feature pruning using Banzhaf power indices. *arXiv preprint arXiv:1711.04992*, 2017.

Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30, 2017.

Maximilian Muschalik, Hubert Baniecki, Fabian Fumagalli, Patrick Kolpaczki, Barbara Hammer, and Eyke Hüllermeier. KernelSHAP-IQ: Weighted least-square optimization for Shapley interactions. *arXiv preprint arXiv:2405.10852*, 2024.

Guillermo Owen. Multilinear extensions of games. *Management Science*, 18(5):P64–P79, 1972.

Giovanni Rabitti and Emanuele Borgonovo. A Shapley-Owen index for interaction quantification. *SIAM/ASA Journal on Uncertainty Quantification*, 7(3):1060–1075, 2019.

Shuai Shao, Ye Tian, Yajie Zhang, and Xingyi Zhang. Evolutionary computation for sparse multiobjective optimization: A survey. *ACM Computing Surveys*, 57(11), 2025.

Lloyd S. Shapley. A value for n-person games. In Harold W. Kuhn and Albert W. Tucker (eds.), *Contributions to the Theory of Games*, volume 2, pp. 307–317. Princeton University Press, 1953.

Lloyd S. Shapley and Martin Shubik. A method for evaluating the distribution of power in a committee system. *American Political Science Review*, 48(3):787–792, 1954.

Eunhye Song, Barry L. Nelson, and Jeremy Staum. Shapley effects for global sensitivity analysis: Theory and computation. *SIAM/ASA Journal on Uncertainty Quantification*, 4:1060–1083, 2016.

Xianfang Song, Yong Zhang, Wanqiu Zhang, Chunlin He, Ying Hu, Jian Wang, and Dunwei Gong. Evolutionary computation for feature selection in classification: A comprehensive survey of solutions, applications and challenges. *Swarm and Evolutionary Computation*, 90:101661, 2024.

Philip D. Straffin. The Shapley-Shubik and Banzhaf power indices as probabilities. In Alvin E. Roth (ed.), *The Shapley Value: Essays in Honor of Lloyd S. Shapley*, pp. 71–81. Cambridge University Press, 1977.

Mukund Sundararajan and Amir Najmi. The many Shapley values for model explanation. In *International Conference on Machine Learning (ICML)*, 2020.

Che-Ping Tsai, Chih-Kuan Yeh, and Pradeep Ravikumar. Faith-Shap: The faithful Shapley interaction index. *Journal of Machine Learning Research*, 24, 2023.

Jarne Verhaeghe, Jonas Van Der Donckt, Femke Ongenae, and Sofie Van Hoecke. Powershap: A power-full Shapley feature selection method. In *Proceedings of the European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML PKDD)*, 2022.

Bing Xue, Mengjie Zhang, Will N. Browne, and Xin Yao. A survey on evolutionary computation approaches to feature selection. *IEEE Transactions on Evolutionary Computation*, 20(4):606–626, 2016.

Chi Zhao, Jiapeng Liu, and Elena Parilina. ShapG: New feature importance method based on the Shapley value. *Engineering Applications of Artificial Intelligence*, 148, 2025.

Chen Zheng and Garvesh Raskutti. MinShap: A modified Shapley value approach for feature selection. *arXiv preprint arXiv:2604.15107*, 2026. Cited as conceptual inspiration; our theoretical development and algorithmic design are independent.

A Appendix

You may include other additional sections here.